

# RESULTS

---

- Validation sets returned 91% accuracy.
- Test sets returned 88% accuracy.
- Changes in min\_df and max\_df led to minimal gains.
- Star ratings were predicted with greater accuracy. This can be used to flag users who gave the wrong amount of stars.
- Deep learning did not initially outperform Logistic Regression, but more tests and reviews deliver better results.
- The helpful rating metric can be applied to any product that counts votes (traditionally thumbs up / thumbs down).
- The next step is to use a similar pipeline to determine helpful reviewers.



# REFERENCES

---

All reports, data wrangling, data analysis, and machine learning jupyter notebooks are on github.

[https://github.com/coreyjwade/Helpful\\_Reviews](https://github.com/coreyjwade/Helpful_Reviews)

Publicly available Amazon datasets.

<http://jmcauley.ucsd.edu/data/amazon/>

My personal website.

[coreyjwade.com](http://coreyjwade.com)